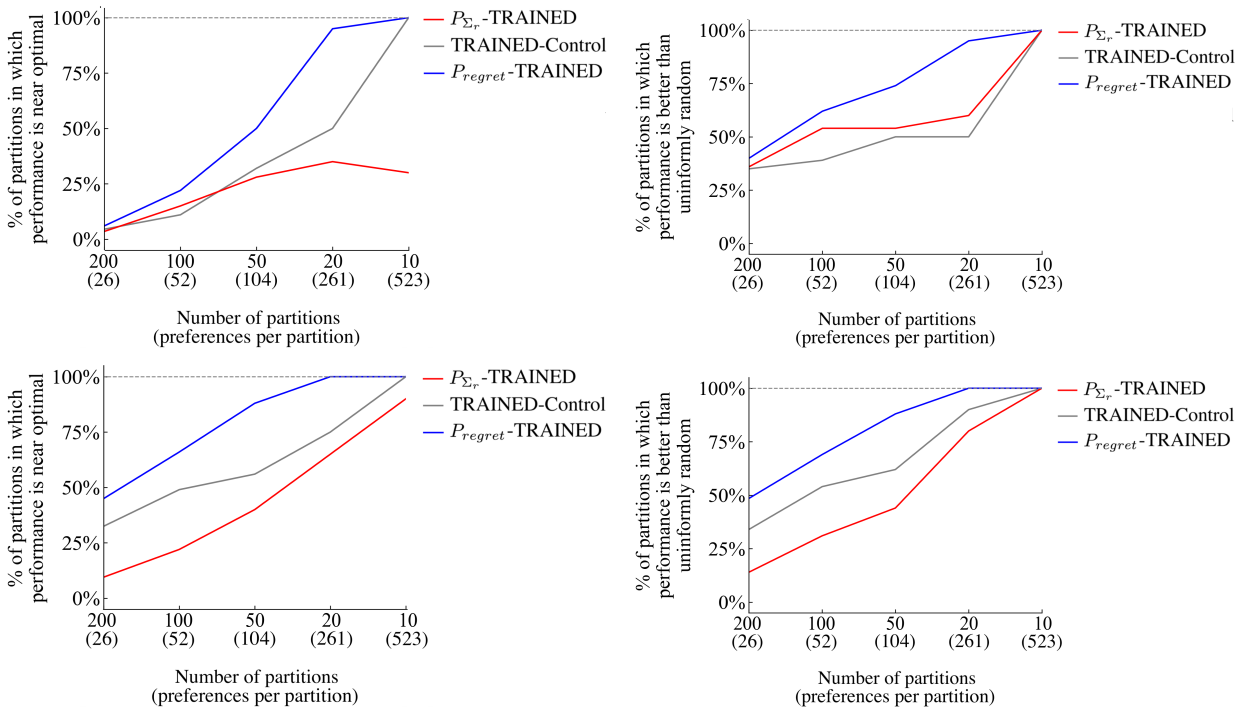

Figure 27: Learning a reward function with the partial return preference model (Top Row) and regret preference model (Bottom Row) from *synthetic* and human preferences. The TRAINED experiment's preference datasets are partitioned following the same methodology as when generating Figure 8. Each preference dataset contains 50 additional segment pairs with synthetic preference labels that aim to compensate for the identifaibility issues of the partial return preference model.